

# Recursion in ARM

## Factorial function

The register R0 is always used to pass the first argument in a function, however since we are using recursion, we are reusing R0 to pass  $n * \text{fact}(n-1)$ .

Thus the compiler moves R0 to R4. Thus contents of R4 is also pushed onto stack

0x5C: 0x0000	10:011	1110110111111111, LRY
x = fact(0);		
0x5E: 0x2000	MOVS	R0, #0

0 is passed via R0

fact:	PUSH	{R4, LR}
0x40: 0xb510		
0x42: 0x0004	MOVS	R4, R0

if(n == 1 || n == 0)

content of R4 is saved onto stack before R0 content is passed to R4

After the recursion call is finished the growing stack shrinks down to the base address which now only contains the argument to be returned (0x78 or 120 which is 5 factorial). This argument is returned via R0:

Registers	Find:	Group:
Name	Value	
R0	0x0000'0078	
R1	0xe000'ed88	
R2	0x0000'0000	
R3	0x0000'0000	
R4	0x0000'0000	
R5	0x0000'0000	
R6	0x0000'0000	
R7	0x0000'0000	
R8	0x0000'0000	
R9	0x0000'0000	
R10	0x0000'0000	
R11	0x0000'0000	
R12	0x0000'0000	
APSR	0x2000'0000	
IPSR	0x0000'0000	
EPSR	0x0100'0000	
PC	0x0000'0062	
SP	0x2000'0fe8	
LR	0x0000'0053	
PRIMASK	0x0000'0000	
BASEPRI	0x0000'0000	
BASEPRI_MAX	0x0000'0000	
FAULTMASK	0x0000'0000	
CONTROL	0x0000'0004	

0x2000'0fc0	00000002
0x2000'0fc4	00000053
0x2000'0fc8	00000003
0x2000'0fcc	00000053
0x2000'0fd0	00000004
0x2000'0fd4	00000053
0x2000'0fd8	00000005
0x2000'0fdc	00000053
0x2000'0fe0	00000000
0x2000'0fe4	00000061
0x2000'0fe8	00000078
0x2000'0fec	00000000
0x2000'0ff0	00000000

0x2000'0fe8 was base address of stack pointer. The LR address is also popped to the PC register to get back to the code just after the first function call.

## The "Stack1" view

Allows to see the entire stack. The bottom address is the beginning address of the stack, and it grows bottom up :

Location	Data	Variable
0x2000'0ff8	0xefef5'eda5	
0x2000'0ffc	0xefef5'eda5	

→ beginning stack address

Location	Data	Variable
0x2000'0fb8	0x1	foo[0]
0x2000'0fbc	0x53	foo[1]
0x2000'0fc0	0x2	foo[2]
0x2000'0fc4	0x53	foo[3]
0x2000'0fc8	0x3	foo[4]
0x2000'0fcc	0x5	foo[5]
0x2000'0fd0	0x4	foo[6]
0x2000'0fd4	0x53	foo[7]
0x2000'0fd8	0x5	foo[8]
0x2000'0fdc	0x53	foo[9]
0x2000'0fe0	0x0	
0x2000'0fe4	0x6d	
0x2000'0fe8	0x0	
0x2000'0fec	0x0	
0x2000'0ff0	0x0	
0x2000'0ff4	0x137	
0x2000'0ff8	0xefef5'eda5	
0x2000'0ffc	0xefef5'eda5	

↑  
Stack grows bottom-up

## Stack Overflow

```
unsigned fact(unsigned n)
{
    unsigned foo[10];
    foo[n] = n;

    if(n == 0u)
    {
        return 1u;
    }
    return foo[n]*fact(n - 1u);
}
```

function being executed. This gets saved onto the stack.

The passed argument is n=5 and hence foo[5] contains 0x5 while all other foo's contain garbage values.

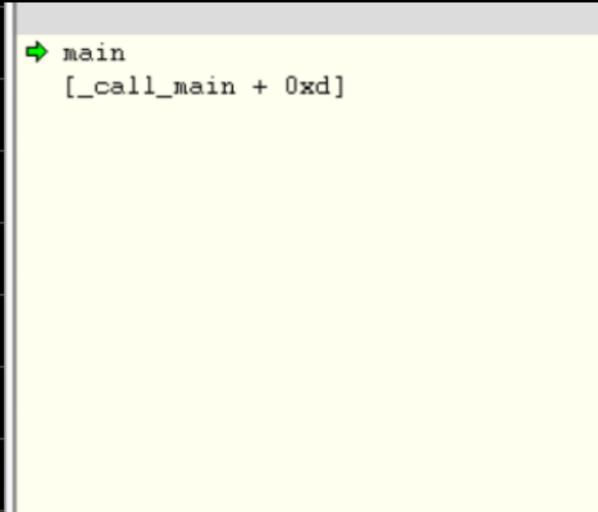
Location	Data	Variable
0x2000'0fb8	0x1	foo[0]
0x2000'0fbc	0x53	foo[1]
0x2000'0fc0	0x2	foo[2]
0x2000'0fc4	0x53	foo[3]
0x2000'0fc8	0x3	foo[4]
0x2000'0fcc	0x5	foo[5]
0x2000'0fd0	0x4	foo[6]
0x2000'0fd4	0x53	foo[7]
0x2000'0fd8	0x5	foo[8]
0x2000'0fdc	0x53	foo[9]
0x2000'0fe0	0x0	
0x2000'0fe4	0x6d	
0x2000'0fe8	0x0	
0x2000'0fec	0x0	
0x2000'0ff0	0x0	
0x2000'0ff4	0x137	
0x2000'0ff8	0xefef5'eda5	
0x2000'0ffc	0xefef5'eda5	

If we 'step into' our code again then we see the stack grow again due to foo and only foo[4] contains correct value which is 0x4. The rest of the foos are again garbage.

0x2000'0f88	0xd004'4f60	foo[0]
0x2000'0f8c	0xf04f'4658	foo[1]
0x2000'0f90	0xf7ff'4160	foo[2]
0x2000'0f94	0x9b0d'fd27	foo[3]
0x2000'0f98	0x4	foo[4]
0x2000'0f9c	0x4650'd059	foo[5]
0x2000'0fa0	0x4160'f04f	foo[6]
0x2000'0fa4	0xf1b8'e053	foo[7]
0x2000'0fa8	0xd003'0fef	foo[8]
0x2000'0fac	0x3fe'f24e	foo[9]
0x2000'0fb0	0x5	
0x2000'0fb4	0x5d	
0x2000'0fb8	0x1	foo[0]
0x2000'0fdc	0x53	foo[1]
0x2000'0fc0	0x2	foo[2]
0x2000'0fc4	0x53	foo[3]
0x2000'0fc8	0x3	foo[4]
0x2000'0fcc	0x5	foo[5]
0x2000'0fd0	0x4	foo[6]
0x2000'0fd4	0x53	foo[7]
0x2000'0fd8	0x5	foo[8]
0x2000'0fdc	0x53	foo[9]
0x2000'0fe0	0x0	
0x2000'0fe4	0x6d	

The stack growing this much is an instance of possible stack overflow. We can demonstrate this further by setting foo[] size to 100 (set a breakpoint at return foo[n]\*n-1 statement).

First we change to the "call stack" view in IAR. This shows all function calls currently nested on stack. "main" itself is called from the stack as shown by [-call-main +0x9], this place contains startup code.



for fact(10) and foo[100], stack overflow occurs after 7 calls to function fact. Here program freezes and the stack pointer is below the valid RAM in ARM (it is at 0xfffffff90 while RAM starts at 0x20000000).

Stack 1			Registers 1		
CSTACK			Find:	Group:	Current CPU Registers
Location	Data	Variable	Name	Value	
0xffff'ff90	-		R9	0x0000'0000	
0xffff'ff94	-		R10	0x0000'0000	
0xffff'ff98	-		R11	0x0000'0000	
0xffff'ff9c	-		R12	0x0000'0000	
0xffff'ffa0	-		APSR	0x2000'0000	
0xffff'ffa4	-		IPSR	0x0000'0003	
0xffff'ffa8	-		EPSR	0x0100'0000	
0xffff'ffac	-		PC	0x0000'011a	
0xffff'ffb0	-		SP	0xffff'ff90	
0xffff'ffb4	-		LR	0xffff'ffe9	
0xffff'ffb8	-		PRIMASK	0x0000'0000	
0xffff'ffbc	-		BASEPRI	0x0000'0000	
0xffff'ffcc	-		BASEPRI_MAX	0x0000'0000	

A look at disassembly tells us that the code is at "BusFault\_Handler" and this is an interrupt by ARM used to stop program flow when non-existent memory. This is an endless loop.

We can change this stack size in IAR by ticking "override default" in config tab in project options.

We can edit the stack (cstack) by clicking the edit button below and going to Stack/heap option.

We can also set heap to 0 since it is a bad idea to use in embedded systems.

### Another example for Stack Overflow

Set foo size to foo[6] and call fact with argument 7 [fact(7)]. Set a breakpoint at fact(7) and single step from there. After the usual push R4,LR to stack command there is a SUB sp,sp,#0x18 performed

```
unsigned fact(unsigned n)
{
    fact:
    0x40: 0xb510      PUSH   {R4, LR}
    0x42: 0xb086      SUB    SP, SP, #0x18
    foo[n] = n;
    0x44: 0x4669      MOV    R1, SP
    0x46: 0xf841 0x0020 STR.W R0, [R1, R0, LSL]
    if(n == 0u)
    0x4a: 0x2800      CMP    R0, #0
    0x4c: 0xd101      BNE.N ??fact_0
    return 1u;
```

This SUB instruction subtracts 0x18 from our stack address and makes stack larger to accommodate for foo[6] array

Location	Data	Variable
0x2000'03e0	0x0	
0x2000'03e4	0x6d	
0x2000'03e8	0x0	
0x2000'03ec	0x0	
0x2000'03f0	0x0	
0x2000'03f4	0x137	
0x2000'03f8	0xfef5'eda5	
0x2000'03fc	0xfef5'eda5	

Name	Value
R12	0x0000'0000
APSR	0x2000'0000
IPSR	0x0000'0000
EPSR	0x0100'0000
PC	0x0000'0042
SP	0x2000'03e0
LR	0x0000'006d
PRIMASK	0x0000'0000

stack pointer at 0x2000'03e0 goes to 0x2000'03c8 to make room for foo[6] (-0x018)

Location	Data	Variable
0x2000'03c8	0x1188	foo[0]
0x2000'03cc	0x4	foo[1]
0x2000'03d0	0x0	foo[2]
0x2000'03d4	0x7c00'0000	foo[3]
0x2000'03d8	0x5901'8400	foo[4]
0x2000'03dc	0xffff	foo[5]
0x2000'03e0	0x0	
0x2000'03e4	0x6d	
0x2000'03e8	0x0	
0x2000'03ec	0x0	
0x2000'03f0	0x0	
0x2000'03f4	0x137	
0x2000'03f8	0xfef5'eda5	

Name	Value
R12	0x0000'0000
APSR	0x2000'0000
IPSR	0x0000'0000
EPSR	0x0100'0000
PC	0x0000'0044
SP	0x2000'03c8
LR	0x0000'006d
PRIMASK	0x0000'0000
BASEPRI	0x0000'0000
BASEPRI_MAX	0x0000'0000
FAULTMASK	0x0000'0000
CONTROL	0x0000'0004
CYCLECOUNTER	0

Next, the MOV instruction moves SP address to R1 and the STR R0,[R1,R0,LSL#2] instruction shifts the argument passed to our function (and still stored in R0) to the R0 register again. LSL is logical shift left

However our index  $n=7$  is beyond the array bound of  $\text{foo}[5]$ . Hence  $\text{foo}[n]=n$  will go two places beyond  $\text{foo}[5]$ 's address ( $0x2000\ 03dc$ ) to  $0x2000\ 02e4$  which contains our link register data. This corrupts LR data and the stack.

To troubleshoot this, add a breakpoint at the return address of function (just before POP instruction):

```
Disassembly
0x42: 0xb086      SUB    SP, SP, #0x18
foo[n] = n;
0x44: 0x4669      MOV    R1, SP
0x46: 0xf841 0x0020 STR.W  R0, [R1, R0, LSI
if(n == 0u)
0x4a: 0x2800      CMP    R0, #0
0x4c: 0xd101      BNE.N ??fact_0
return 1u;
0x4e: 0x2001      MOVS   R0, #1
0x50: 0xe006      B.N    ??fact_1
return foo[n]*fact(n - 1u);
??fact_0:
0x52: 0xf851 0x4020 LDR.W  R4, [R1, R0, LSI
0x56: 0x1e40      SUBS   R0, R0, #1
0x58: 0xf7ff 0xffff2 BL     fact
0x5c: 0x4344      MULS   R4, R0, R4
0x5e: 0x0020      MOVS   R0, R4
??fact_1:
0x60: 0xb006      ADD    SP, SP, #0x18
0x62: 0xbd10      POP    {R4, PC}
int main()
{
    main:
0x64: 0xb51c      PUSH   {R2-R4, LR}
```

The Add  $SP, SP, \#0x18$  adds  $0x18$  to stack address and reduces its size. The POP instruction restores the saved (and now corrupted) link register into the PC. Recall this place hold data  $0x07$  due to stack corruption when  $\text{foo}[7]$  crossed foo array bounds into the place where LR was saved on the stack. This causes wrong program flow and puts the CPU into a cycle. Throughout all this stack is not properly ejected, and thus grows till there is stack overflow.

at this breakpoint all calls are complete and will recursively unwind from here.

Stack 1

Location	Data	Variable
0x2000'02e8	0x0	foo[0]
0x2000'02ec	0xb001'dbf8	foo[1]
0x2000'02f0	0xbf00'4770	foo[2]
0x2000'02f4	0xf'4240	foo[3]
0x2000'02f8	0x5188'f64e	foo[4]
0x2000'02fc	0x100'f2ce	foo[5]
0x2000'0300	0x1	
0x2000'0304	0x5d	
0x2000'0308	0x8f4f'f3bf	foo[0]
0x2000'030c	0x1	foo[1]
0x2000'0310	0x7001'f04f	foo[2]
0x2000'0314	0xa10'eee1	foo[3]
0x2000'0318	0xe7fe'4770	foo[4]
0x2000'031c	0xf80d'f000	foo[5]
0x2000'0320	0x2	
0x2000'0324	0x5d	
0x2000'0328	0x8000'f3af	foo[0]
0x2000'032c	0xf3af'2000	foo[1]
0x2000'0330	0x2	foo[2]
0x2000'0334	0xf000'ff97	foo[3]
0x2000'0338	0x2001'f802	foo[4]
0x2000'033c	0xf000'4770	foo[5]
0x2000'0340	0x3	
0x2000'0344	0x5d	

Disassembly

```
0x106: 0x6008      STR    R0, [R1]
0x108: 0xf3bf 0x8f4f DSB
0x10c: 0xf3bf 0x8f6f ISB
0x110: 0xf04f 0x7001 MOV.W R0, #33816576
0x114: 0xeeee 0xa0a10 VMSR   FPSCR, R0
0x118: 0x4770      BX     LR
    BusFault_Handler:
    DebugMon_Handler:
    HardFault_Handler:
    MemManage_Handler:
    NMI_Handler... +5 symbols not displayed
0x11a: 0xe7fe      B.N    BusFault_Handler
    ?main:
    __main:
0x11c: 0xf000 0xf80d BL     __low_level_init
0x120: 0x2800      CMP    R0, #0
0x122: 0xd001      BEQ.N _call_main
0x124: 0xf3af 0x8000 NOP.W
    _call_main:
0x128: 0xf3af 0x8000 NOP.W
0x12c: 0x2000      MOVS   R0, #0
0x12e: 0xf3af 0x8000 NOP.W
0x132: 0xf7ff 0xff97 BL     main
0x136: 0xf000 0xf802 BL     exit
```

```

void swap(volatile int *x, volatile int *y)
{
    int tmp;
    tmp = *x;
    *x = *y;
    *y = tmp;
}

```

```
int main()
```

```
{
```

```
volatile int x;
```

```
volatile int y;
```

```
x = 5u;
```

```
y = 2u;
```

```
swap(&x, &y);
```

```
}
```

Since swap is using pointers, all changes in x and y of swap mean x and y of main is getting altered directly

```

int main()
{
    main:
    0x4c: 0xb51c      PUSH   {R2-R4, LR}
    x = 5u;
    0x4e: 0x2005      MOVS   R0, #5
    0x50: 0x9001      STR    R0, [SP, #0x4]
    y = 2u;
    0x52: 0x2002      MOVS   R0, #2
    0x54: 0x9000      STR    R0, [SP]
    swap(&x, &y);
    0x56: 0x4669      MOV    R1, SP
    0x58: 0xa801      ADD    R0, SP, #0x4
    0x5a: 0xf7ff 0xffff1 BL    swap

```

x and y are stored in the stack using STR function.  
X is stored at address:  
SP address + 0x04 = 0x200003ec  
and y is stored in top SP address: 0x200003e8

Stack 1			
	CSTACK		X
Location	Data	Variable	Value
0x2000'03e8	0x2	y	2
0x2000'03ec	0x5	x	5
0x2000'03f0	0x0		
0x2000'03f4	0x127		
0x2000'03f8	0xfef5'eda5		
0x2000'03fc	0xfef5'eda5		

As per ARM standard, address of R0 and R1 in stack is copied to R1 and R0 registers just before branch (BL) instruction is carried out:

Name	Value	Disassembly
R0	0x2000'03ec	*y = **x;
R1	0x2000'03e8	0x46: 0x6800 LDR R0, [R0]
R2	0x0000'0000	0x48: 0x6008 STR R0, [R1]
R3	0x0000'0000	}
R4	0x0000'0000	0x4a: 0x4770 BX LR
R5	0x0000'0000	int main()
R6	0x0000'0000	{
R7	0x0000'0000	main:
R8	0x0000'0000	0x4c: 0xb51c PUSH {R2-R4, LR}
R9	0x0000'0000	x = 5u;
R10	0x0000'0000	0x4e: 0x2005 MOVS R0, #5
R11	0x0000'0000	0x50: 0x9001 STR R0, [SP, #0x4]
R12	0x0000'0000	y = 2u;
		0x52: 0x2002 MOVS R0, #2
		0x54: 0x9000 STR R0, [SP]
		swap(&x, &y);
		0x56: 0x4669 MOV R1, SP
		0x58: 0xa801 ADD R0, SP, #0x4
		0x5a: 0xf7ff 0xffff1 BL swap

The next set of instructions show how swap function changes the value of x and y themselves via utilising pointers :

swap:		
0x40: 0x6802	LDR	R2, [R0]
*x = *y;		
0x42: 0x680b	LDR	R3, [R1]
0x44: 0x6003	STR	R3, [R0]
*y = tmp;		
0x46: 0x600a	STR	R2, [R1]
}		

1st LDR moves content of address stored in R0 (value of x) to R2

2nd LDR moves content of address stored in R1 (value of y) to R3

3rd STR stores R3 contents at address stored in R0

4th STR stores R2 contents in address stored in R1. This causes x and y swap.

\* Imp for returning local variables in function

Note → stack destroys previous values of local variables from a function, anytime another function call is initiated. If the variables and their values need to be used after returning to main, it is best to declare the function local variables as static